

ExampleB

March 26, 2021

1 Project enda : Example B

If you haven't already, read Example A first, it is not long. Download `example_b.zip` and run this notebook in the correct python environment.

Install all the required packages in your python virtualenv:

```
pip install numexpr bottleneck pandas enda jupyter h2o scikit-learn statsmodels matplotlib joblib  
pip install jours-feries-france vacances-scolaires-france Unidecode # used for feature engineering
```

In this example we will go more in depth, with realistic data and more historical data (~4-5 years). This example is divided in 7 parts: 1. Read and prepare data, check for missing values and gaps 2. Visualize data 3. Feature engineering : datetime and calendar features 4. Portfolio forecast & basic prediction 5. Benchmark with simple evaluation 6. Benchmark with Backtesting 7. Make the prediction

We set ourselves in a setup as if we were **exactly on 2020-11-30**. We want to predict the total consumption of customers for the next few days starting 2020-12-01 at a 30min time-step. We have: - our customer contracts until 2020-11-30 included. - historical load data from 2015-01-01 until 2020-11-15 included. There is a ~15 day time-gap between the last moment for which we have an actual load measure and 'today' (2020-11-30). - weather forecast until 2020-12-11 (11 days). - our TSO's network load forecast until 2020-12-7 (7 days).

In here (example B), we will put all our customers in only 1 group and forecast the next 7 days. We will first construct the dataset and the forecast input data and test it with a basic linear regressor. We will then try various algorithms and compare them. Finally we will give an example of backtesting on the data.

```
[1]: import enda  
import pandas as pd  
import os
```

1.1 1. Read and prepare data, check for missing values and gaps

```
[2]: # Replace this with the path to your example_b directory.  
# You should have ExampleB.ipynb opened in jupyter, so you can run each step  
DIR = '/Users/emmanuel.charon/Documents/CodeProjects/enercoop/enda/data/  
→example_b'
```

```

[3]: # Get the 30min time-step data just like in Example A
# (columns are a bit different and there is more data)
# Here we consider all customers in one big group.
def read_data():
    contracts = enda.Contracts.read_contracts_from_file(os.path.join(DIR,
↪ "contracts.csv"))
    contracts["contracts_count"] = 1
    portfolio_by_day = enda.Contracts.compute_portfolio_by_day(
        contracts,
        columns_to_sum = ["contracts_count", "kva"],
        date_start_col="date_start",
        date_end_exclusive_col="date_end_exclusive",
    )
    portfolio = enda.TimeSeries.interpolate_daily_to_sub_daily_data(
        portfolio_by_day,
        freq='30min',
        tz='Europe/Paris'
    )

    historic_load_measured = pd.read_csv(os.path.join(DIR,
↪ "historic_load_measured.csv"))
    weather_and_tso_forecasts = pd.read_csv(os.path.join(DIR,
↪ "weather_and_tso_forecasts.csv"))
    # correctly format 'time' as a pandas.DatetimeIndex of dtype: datetime[ns,
↪ tzinfo]
    for df in [historic_load_measured, weather_and_tso_forecasts]:
        df['time'] = pd.to_datetime(df['time'])
        df['time'] = enda.TimeSeries.align_timezone(df['time'], tzinfo =
↪ 'Europe/Paris')
        df.set_index('time', inplace=True)

    # keep only where both loads are known
    historic_load_measured = historic_load_measured.dropna()
    historic_load_measured["load_kw"] =
↪ historic_load_measured["smart_metered_kw"] + historic_load_measured["slp_kw"]
    # keep only the full load
    historic_load_measured = historic_load_measured[["load_kw"]]

    return contracts, portfolio, historic_load_measured,
↪ weather_and_tso_forecasts

[4]: contracts, portfolio, historic_load_measured, weather_and_tso_forecasts =
↪ read_data()
# remove data where tso is not available

```

```
weather_and_tso_forecasts = weather_and_tso_forecasts.  
↳dropna(subset=["tso_forecast_load_mw"])
```

```
[5]: contracts
```

```
[5]:
```

	date_start	date_end_exclusive	kva	meter_reading_type	contracts_count
0	2006-08-09	NaT	12.0	PROFILE	1
1	2006-09-01	2006-11-23	6.0	PROFILE	1
2	2006-09-01	2007-11-01	3.0	PROFILE	1
3	2006-09-01	2007-12-19	12.0	PROFILE	1
4	2006-09-01	2008-06-28	12.0	PROFILE	1
...
162598	2020-11-30	NaT	6.0	PROFILE	1
162599	2020-11-30	NaT	6.0	PROFILE	1
162600	2020-11-30	NaT	6.0	PROFILE	1
162601	2020-11-30	NaT	6.0	PROFILE	1
162602	2020-11-30	NaT	6.0	PROFILE	1

[162603 rows x 5 columns]

```
[6]: portfolio
```

```
[6]:
```

	contracts_count	kva
time		
2006-08-09 00:00:00+02:00	1.0	12.0
2006-08-09 00:30:00+02:00	1.0	12.0
2006-08-09 01:00:00+02:00	1.0	12.0
2006-08-09 01:30:00+02:00	1.0	12.0
2006-08-09 02:00:00+02:00	1.0	12.0
...
2020-11-30 21:30:00+01:00	96134.0	820005.7
2020-11-30 22:00:00+01:00	96134.0	820005.7
2020-11-30 22:30:00+01:00	96134.0	820005.7
2020-11-30 23:00:00+01:00	96134.0	820005.7
2020-11-30 23:30:00+01:00	96134.0	820005.7

[250946 rows x 2 columns]

```
[7]: historic_load_measured
```

```
[7]:
```

	load_kw
time	
2015-01-01 00:00:00+01:00	2490.925806
2015-01-01 00:30:00+01:00	2412.623113
2015-01-01 01:00:00+01:00	2365.611276
2015-01-01 01:30:00+01:00	2336.141065
2015-01-01 02:00:00+01:00	2300.935642

```

...
2020-11-15 21:30:00+01:00 7657.293444
2020-11-15 22:00:00+01:00 7317.540759
2020-11-15 22:30:00+01:00 7580.051439
2020-11-15 23:00:00+01:00 7496.273993
2020-11-15 23:30:00+01:00 7376.005701

```

[97198 rows x 1 columns]

```

[8]: # t_weighted is the average french temperature weighted by population density
      # t_smooth is a smoothing computed over t_weighted to take into account
      ↳ building calorific inertia
      # (t_smooth is computed out of enda here)
      weather_and_tso_forecasts

```

```

[8]:
      tso_forecast_load_mw  t_weighted  t_smooth
time
2015-01-01 00:00:00+01:00      72900.0      -0.41      1.17
2015-01-01 00:30:00+01:00      71600.0      -0.48      1.17
2015-01-01 01:00:00+01:00      69900.0      -0.55      1.15
2015-01-01 01:30:00+01:00      70600.0      -0.66      1.14
2015-01-01 02:00:00+01:00      70500.0      -0.78      1.11
...
2020-12-07 21:30:00+01:00      68400.0      4.20      4.13
2020-12-07 22:00:00+01:00      66900.0      4.12      4.10
2020-12-07 22:30:00+01:00      67600.0      4.03      4.08
2020-12-07 23:00:00+01:00      70200.0      3.94      4.07
2020-12-07 23:30:00+01:00      69600.0      3.94      4.07

```

[104064 rows x 3 columns]

```

[9]: # lets create the train set with historical data
      historic = pd.merge(
          portfolio,
          historic_load_measured, # here we select only the load of the desired group
          how='inner', left_index=True, right_index=True
      )

      historic = pd.merge(
          historic,
          weather_and_tso_forecasts,
          how='inner', left_index=True, right_index=True
      )

```

```

[10]: historic

```

```
[10]:
```

	contracts_count	kva	load_kw \
time			
2015-01-01 00:00:00+01:00	21261.0	167416.4	2490.925806
2015-01-01 00:30:00+01:00	21261.0	167416.4	2412.623113
2015-01-01 01:00:00+01:00	21261.0	167416.4	2365.611276
2015-01-01 01:30:00+01:00	21261.0	167416.4	2336.141065
2015-01-01 02:00:00+01:00	21261.0	167416.4	2300.935642
...
2020-11-15 21:30:00+01:00	95475.0	813328.8	7657.293444
2020-11-15 22:00:00+01:00	95475.0	813328.8	7317.540759
2020-11-15 22:30:00+01:00	95475.0	813328.8	7580.051439
2020-11-15 23:00:00+01:00	95475.0	813328.8	7496.273993
2020-11-15 23:30:00+01:00	95475.0	813328.8	7376.005701

	tso_forecast_load_mw	t_weighted	t_smooth
time			
2015-01-01 00:00:00+01:00	72900.0	-0.41	1.17
2015-01-01 00:30:00+01:00	71600.0	-0.48	1.17
2015-01-01 01:00:00+01:00	69900.0	-0.55	1.15
2015-01-01 01:30:00+01:00	70600.0	-0.66	1.14
2015-01-01 02:00:00+01:00	70500.0	-0.78	1.11
...
2020-11-15 21:30:00+01:00	46200.0	12.05	12.01
2020-11-15 22:00:00+01:00	45200.0	11.92	11.97
2020-11-15 22:30:00+01:00	46400.0	11.84	11.96
2020-11-15 23:00:00+01:00	48600.0	11.75	11.94
2020-11-15 23:30:00+01:00	49400.0	11.64	11.92

[97198 rows x 6 columns]

```
[11]: # check that there is no NaN value
historic.isna().sum()
```

```
[11]: contracts_count    0
kva                      0
load_kw                 0
tso_forecast_load_mw    0
t_weighted              0
t_smooth                0
dtype: int64
```

```
[12]: # note that the type of the index is precise
historic.index.dtype, type(historic.index)
```

```
[12]: (datetime64[ns, Europe/Paris], pandas.core.indexes.datetimes.DatetimeIndex)
```

```
[13]: # check missing data in the timeseries (based on the time index only)
freq, missing_periods, extra_points = enda.TimeSeries.
    ↪ find_missing_and_extra_periods(
        dti=historic.index,
        expected_freq = '30min',
        expected_start_datetime = pd.to_datetime('2015-01-01 00:00:00+01:00').
    ↪ astimezone('Europe/Paris'),
        expected_end_datetime = pd.to_datetime('2020-11-30 23:30:00+01:00').
    ↪ astimezone('Europe/Paris')
    )
for missing_period in missing_periods:
    print("Missing data from {} to {}".format(missing_period[0],
    ↪ missing_period[1]))
if len(extra_points) > 0 :
    print("Extra points found: {}".format(extra_points))
```

Missing data from 2015-09-01 00:00:00+02:00 to 2015-11-30 23:30:00+01:00.
 Missing data from 2018-06-01 00:00:00+02:00 to 2018-06-30 23:30:00+02:00.
 Missing data from 2020-11-16 00:00:00+01:00 to 2020-11-30 23:30:00+01:00.

We expected the missing data from 2020-11-16 to 2020-11-30, but not from the rest.

```
[14]: # Zoom on a daylight savings time change to double-check that it was handled
    ↪ correctly
historic[(historic.index >= '2019-10-27 01:00:00+02:00') & (historic.index <
    ↪ '2019-10-27 03:30:00+01:00')]
```

```
[14]:
```

	contracts_count	kva	load_kw \
time			
2019-10-27 01:00:00+02:00	84131.0	716816.4	5179.955556
2019-10-27 01:30:00+02:00	84131.0	716816.4	5087.111111
2019-10-27 02:00:00+02:00	84131.0	716816.4	4898.400000
2019-10-27 02:30:00+02:00	84131.0	716816.4	4616.533333
2019-10-27 02:00:00+01:00	84131.0	716816.4	4259.822222
2019-10-27 02:30:00+01:00	84131.0	716816.4	4208.888889
2019-10-27 03:00:00+01:00	84131.0	716816.4	4137.955556


	tso_forecast_load_mw	t_weighted	t_smooth
time			
2019-10-27 01:00:00+02:00	41300.0	13.65	13.49
2019-10-27 01:30:00+02:00	40700.0	13.52	13.47
2019-10-27 02:00:00+02:00	36700.0	13.40	13.46
2019-10-27 02:30:00+02:00	36700.0	13.26	13.44
2019-10-27 02:00:00+01:00	36700.0	13.12	13.42
2019-10-27 02:30:00+01:00	36700.0	12.91	13.39
2019-10-27 03:00:00+01:00	36700.0	12.70	13.37

1.2 2. Visualize data

In order to visualise using pandas, we use the matplotlib backend.

```
[15]: # Show full data set
historic.plot(figsize=(20, 30), subplots=True)


[15]: array([<AxesSubplot:xlabel='time'>, <AxesSubplot:xlabel='time'>,
        <AxesSubplot:xlabel='time'>, <AxesSubplot:xlabel='time'>,
        <AxesSubplot:xlabel='time'>, <AxesSubplot:xlabel='time'>],
        dtype=object)
```



```
[16]: # Show recent data
historic[-1000:].plot(figsize=(20, 30), subplots=True)
```



```
[16]: array([<AxesSubplot:xlabel='time'>, <AxesSubplot:xlabel='time'>,  
            <AxesSubplot:xlabel='time'>, <AxesSubplot:xlabel='time'>,  
            <AxesSubplot:xlabel='time'>, <AxesSubplot:xlabel='time'>],  
          dtype=object)
```



Don't hesitate to add your own visualisations!

```
[ ]:
```

1.3 3. Feature engineering

Before we train, we will add some features based on the `datetime`, and some calendar features related to national holidays or school holidays.

We use some packages for the holidays, which are used in `enda.feature_engineering.calendar`:

```
pip install jours-feries-france vacances-scolaires-france Unidecode
```

```
[17]: import enda.feature_engineering.calendar
```

```
[18]: # define the features we want to add before training/predicting
def featurize(df):
    # put datetime features to capture the data frequencies: daily, weekly and
    ↪ yearly periods.
    df = enda.DatetimeFeature.split_datetime(
        df, split_list = ['minuteofday', 'dayofweek', 'month']
    )
    df = enda.DatetimeFeature.encode_cyclic_datetime_index(
        df, split_list = ['minuteofday', 'dayofweek', 'dayofyear']
    )

    # add features about national holidays and school holidays (French holidays
    ↪ here)
    special_days = enda.feature_engineering.calendar.Calendar().
    ↪ get_french_special_days()
    df = pd.merge(
        df, special_days,
        how='left', left_index=True, right_index=True
    )
    return df
```

```
[19]: full_train_set = featurize(historic)
```

```
[20]: full_train_set
```

```
[20]:
```

	contracts_count	kva	load_kw \
time			
2015-01-01 00:00:00+01:00	21261.0	167416.4	2490.925806
2015-01-01 00:30:00+01:00	21261.0	167416.4	2412.623113
2015-01-01 01:00:00+01:00	21261.0	167416.4	2365.611276
2015-01-01 01:30:00+01:00	21261.0	167416.4	2336.141065

2015-01-01 02:00:00+01:00	21261.0	167416.4	2300.935642
...
2020-11-15 21:30:00+01:00	95475.0	813328.8	7657.293444
2020-11-15 22:00:00+01:00	95475.0	813328.8	7317.540759
2020-11-15 22:30:00+01:00	95475.0	813328.8	7580.051439
2020-11-15 23:00:00+01:00	95475.0	813328.8	7496.273993
2020-11-15 23:30:00+01:00	95475.0	813328.8	7376.005701

time	tso_forecast_load_mw	t_weighted	t_smooth \
2015-01-01 00:00:00+01:00	72900.0	-0.41	1.17
2015-01-01 00:30:00+01:00	71600.0	-0.48	1.17
2015-01-01 01:00:00+01:00	69900.0	-0.55	1.15
2015-01-01 01:30:00+01:00	70600.0	-0.66	1.14
2015-01-01 02:00:00+01:00	70500.0	-0.78	1.11
...
2020-11-15 21:30:00+01:00	46200.0	12.05	12.01
2020-11-15 22:00:00+01:00	45200.0	11.92	11.97
2020-11-15 22:30:00+01:00	46400.0	11.84	11.96
2020-11-15 23:00:00+01:00	48600.0	11.75	11.94
2020-11-15 23:30:00+01:00	49400.0	11.64	11.92

time	minuteofday	dayofweek	month	minuteofday_cos \
2015-01-01 00:00:00+01:00	0	3	1	1.000000
2015-01-01 00:30:00+01:00	30	3	1	0.991445
2015-01-01 01:00:00+01:00	60	3	1	0.965926
2015-01-01 01:30:00+01:00	90	3	1	0.923880
2015-01-01 02:00:00+01:00	120	3	1	0.866025
...
2020-11-15 21:30:00+01:00	1290	6	11	0.793353
2020-11-15 22:00:00+01:00	1320	6	11	0.866025
2020-11-15 22:30:00+01:00	1350	6	11	0.923880
2020-11-15 23:00:00+01:00	1380	6	11	0.965926
2020-11-15 23:30:00+01:00	1410	6	11	0.991445

time	minuteofday_sin	dayofweek_cos	dayofweek_sin \
2015-01-01 00:00:00+01:00	0.000000	-0.900969	0.433884
2015-01-01 00:30:00+01:00	0.130526	-0.900969	0.433884
2015-01-01 01:00:00+01:00	0.258819	-0.900969	0.433884
2015-01-01 01:30:00+01:00	0.382683	-0.900969	0.433884
2015-01-01 02:00:00+01:00	0.500000	-0.900969	0.433884
...
2020-11-15 21:30:00+01:00	-0.608761	0.623490	-0.781831
2020-11-15 22:00:00+01:00	-0.500000	0.623490	-0.781831
2020-11-15 22:30:00+01:00	-0.382683	0.623490	-0.781831

2020-11-15 23:00:00+01:00	-0.258819	0.623490	-0.781831
2020-11-15 23:30:00+01:00	-0.130526	0.623490	-0.781831

	dayofyear_cos	dayofyear_sin	lockdown \
time			
2015-01-01 00:00:00+01:00	1.000000	0.000000	0.0
2015-01-01 00:30:00+01:00	1.000000	0.000000	0.0
2015-01-01 01:00:00+01:00	1.000000	0.000000	0.0
2015-01-01 01:30:00+01:00	1.000000	0.000000	0.0
2015-01-01 02:00:00+01:00	1.000000	0.000000	0.0
...
2020-11-15 21:30:00+01:00	0.691771	-0.722117	0.0
2020-11-15 22:00:00+01:00	0.691771	-0.722117	0.0
2020-11-15 22:30:00+01:00	0.691771	-0.722117	0.0
2020-11-15 23:00:00+01:00	0.691771	-0.722117	0.0
2020-11-15 23:30:00+01:00	0.691771	-0.722117	0.0

	public_holiday	nb_school_areas_off \
time		
2015-01-01 00:00:00+01:00	1.0	3.0
2015-01-01 00:30:00+01:00	1.0	3.0
2015-01-01 01:00:00+01:00	1.0	3.0
2015-01-01 01:30:00+01:00	1.0	3.0
2015-01-01 02:00:00+01:00	1.0	3.0
...
2020-11-15 21:30:00+01:00	0.0	0.0
2020-11-15 22:00:00+01:00	0.0	0.0
2020-11-15 22:30:00+01:00	0.0	0.0
2020-11-15 23:00:00+01:00	0.0	0.0
2020-11-15 23:30:00+01:00	0.0	0.0

	extra_long_weekend
time	
2015-01-01 00:00:00+01:00	0.0
2015-01-01 00:30:00+01:00	0.0
2015-01-01 01:00:00+01:00	0.0
2015-01-01 01:30:00+01:00	0.0
2015-01-01 02:00:00+01:00	0.0
...	...
2020-11-15 21:30:00+01:00	0.0
2020-11-15 22:00:00+01:00	0.0
2020-11-15 22:30:00+01:00	0.0
2020-11-15 23:00:00+01:00	0.0
2020-11-15 23:30:00+01:00	0.0

[97198 rows x 19 columns]

```
[21]: # train a basic SKLearnLinearRegression
from enda.ml_backends.sklearn_estimator import SklearnEstimator
from sklearn.linear_model import LinearRegression

lin_reg = SklearnEstimator(LinearRegression())
lin_reg.train(full_train_set, target_col='load_kw')
```

1.4 4. Portfolio forecast & basic prediction

We need an estimate of our portfolio in the next few days, the tso_load and weather forecasts.

In order to get our portfolio in the next few days, here we will just consider the latest trends in our portfolio.

In another setup, you might want to connect to your sales software or ERP and take into account contracts that will end or start soon.

We will use `enda.Contracts.forecast_using_trend` which requires the `statsmodel` package :

```
pip install statsmodels
```

```
[22]: # we will forecast the portfolio using holt method
forecast_portfolio = enda.Contracts.forecast_using_trend(
    portfolio_df=portfolio,
    start_forecast_date=pd.to_datetime("2020-12-01 00:00:00+01:00"),
    nb_days=7,
    past_days=150 # only use recent portfolio trend to forecast the next few
    ↪ days
)
forecast_portfolio
```

```
/Users/emmanuel.charon/Documents/CodeProjects/enercoop/enda/venv/lib/python3.7/s
ite-packages/statsmodels/tsa/holtwinters/model.py:922: ConvergenceWarning:
Optimization failed to converge. Check mle_retvals.
ConvergenceWarning,
```

```
[22]:
```

	contracts_count	kva
time		
2020-12-01 00:00:00+01:00	96134.6	820008.8
2020-12-01 00:30:00+01:00	96135.3	820011.8
2020-12-01 01:00:00+01:00	96135.9	820014.9
2020-12-01 01:30:00+01:00	96136.5	820017.9
2020-12-01 02:00:00+01:00	96137.1	820021.0
...
2020-12-07 21:30:00+01:00	96341.2	821020.6
2020-12-07 22:00:00+01:00	96341.8	821023.7
2020-12-07 22:30:00+01:00	96342.4	821026.7
2020-12-07 23:00:00+01:00	96343.1	821029.8
2020-12-07 23:30:00+01:00	96343.7	821032.8

[336 rows x 2 columns]

```
[23]: # add weather_and_tso_forecasts
forecast_input_data = pd.merge(
    forecast_portfolio,
    weather_and_tso_forecasts,
    how='inner', left_index=True, right_index=True
)
# add feature engineering
forecast_input_data = featurize(forecast_input_data)
forecast_input_data
```

```
[23]:
```

	contracts_count	kva	tso_forecast_load_mw	\
time				
2020-12-01 00:00:00+01:00	96134.6	820008.8	66100.0	
2020-12-01 00:30:00+01:00	96135.3	820011.8	64200.0	
2020-12-01 01:00:00+01:00	96135.9	820014.9	61900.0	
2020-12-01 01:30:00+01:00	96136.5	820017.9	62800.0	
2020-12-01 02:00:00+01:00	96137.1	820021.0	62300.0	
...	
2020-12-07 21:30:00+01:00	96341.2	821020.6	68400.0	
2020-12-07 22:00:00+01:00	96341.8	821023.7	66900.0	
2020-12-07 22:30:00+01:00	96342.4	821026.7	67600.0	
2020-12-07 23:00:00+01:00	96343.1	821029.8	70200.0	
2020-12-07 23:30:00+01:00	96343.7	821032.8	69600.0	

	t_weighted	t_smooth	minuteofday	dayofweek	\
time					
2020-12-01 00:00:00+01:00	4.69	5.08	0	1	
2020-12-01 00:30:00+01:00	4.82	5.10	30	1	
2020-12-01 01:00:00+01:00	4.96	5.12	60	1	
2020-12-01 01:30:00+01:00	5.04	5.13	90	1	
2020-12-01 02:00:00+01:00	5.13	5.14	120	1	
...	
2020-12-07 21:30:00+01:00	4.20	4.13	1290	0	
2020-12-07 22:00:00+01:00	4.12	4.10	1320	0	
2020-12-07 22:30:00+01:00	4.03	4.08	1350	0	
2020-12-07 23:00:00+01:00	3.94	4.07	1380	0	
2020-12-07 23:30:00+01:00	3.94	4.07	1410	0	

	month	minuteofday_cos	minuteofday_sin	\
time				
2020-12-01 00:00:00+01:00	12	1.000000	0.000000	
2020-12-01 00:30:00+01:00	12	0.991445	0.130526	
2020-12-01 01:00:00+01:00	12	0.965926	0.258819	
2020-12-01 01:30:00+01:00	12	0.923880	0.382683	

2020-12-01 02:00:00+01:00	12	0.866025	0.500000
...
2020-12-07 21:30:00+01:00	12	0.793353	-0.608761
2020-12-07 22:00:00+01:00	12	0.866025	-0.500000
2020-12-07 22:30:00+01:00	12	0.923880	-0.382683
2020-12-07 23:00:00+01:00	12	0.965926	-0.258819
2020-12-07 23:30:00+01:00	12	0.991445	-0.130526

	dayofweek_cos	dayofweek_sin	dayofyear_cos \
time			
2020-12-01 00:00:00+01:00	0.62349	0.781831	0.861702
2020-12-01 00:30:00+01:00	0.62349	0.781831	0.861702
2020-12-01 01:00:00+01:00	0.62349	0.781831	0.861702
2020-12-01 01:30:00+01:00	0.62349	0.781831	0.861702
2020-12-01 02:00:00+01:00	0.62349	0.781831	0.861702
...
2020-12-07 21:30:00+01:00	1.00000	0.000000	0.909308
2020-12-07 22:00:00+01:00	1.00000	0.000000	0.909308
2020-12-07 22:30:00+01:00	1.00000	0.000000	0.909308
2020-12-07 23:00:00+01:00	1.00000	0.000000	0.909308
2020-12-07 23:30:00+01:00	1.00000	0.000000	0.909308

	dayofyear_sin	lockdown	public_holiday \
time			
2020-12-01 00:00:00+01:00	-0.507415	0.0	0.0
2020-12-01 00:30:00+01:00	-0.507415	0.0	0.0
2020-12-01 01:00:00+01:00	-0.507415	0.0	0.0
2020-12-01 01:30:00+01:00	-0.507415	0.0	0.0
2020-12-01 02:00:00+01:00	-0.507415	0.0	0.0
...
2020-12-07 21:30:00+01:00	-0.416125	0.0	0.0
2020-12-07 22:00:00+01:00	-0.416125	0.0	0.0
2020-12-07 22:30:00+01:00	-0.416125	0.0	0.0
2020-12-07 23:00:00+01:00	-0.416125	0.0	0.0
2020-12-07 23:30:00+01:00	-0.416125	0.0	0.0

	nb_school_areas_off	extra_long_weekend
time		
2020-12-01 00:00:00+01:00	0.0	0.0
2020-12-01 00:30:00+01:00	0.0	0.0
2020-12-01 01:00:00+01:00	0.0	0.0
2020-12-01 01:30:00+01:00	0.0	0.0
2020-12-01 02:00:00+01:00	0.0	0.0
...
2020-12-07 21:30:00+01:00	0.0	0.0
2020-12-07 22:00:00+01:00	0.0	0.0
2020-12-07 22:30:00+01:00	0.0	0.0


2020-12-07 23:00:00+01:00	0.0	0.0
2020-12-07 23:30:00+01:00	0.0	0.0

[336 rows x 18 columns]

```
[24]: # do the prediction
lin_reg_prediction = lin_reg.predict(forecast_input_data, target_col="load_kw")
```

```
[25]: # visualize recent load along with our forecast; remember we don't have recent_
      ↪ actual load so there is a time-gap.
to_plot = pd.merge(
    historic["load_kw"][-1000:].to_frame("historic_kw"),
    lin_reg_prediction.rename(columns={"load_kw": "forecast_kw"}),
    how='outer', left_index=True, right_index=True
)
to_plot.plot(ylim=0, figsize=(16, 4))
```

[25]: <AxesSubplot:xlabel='time'>



1.5 5. Benchmark with simple evaluation

The previous forecast based on linear regression is very limited. Let's try and use a better algorithm !

We will define some algorithms using `scikit-learn` as a machine learning backend and others using `h2o`.

For that we need the `h2o` package:

```
pip install h2o
```

```
[26]: # here we do a benchmark, we want to compare with actual data,
      # lets say from 2020-11-01 to 2020-11-15
benchmark_train = full_train_set[full_train_set.index < '2020-11-01']
benchmark_test = full_train_set[full_train_set.index >= '2020-11-01']

benchmark = benchmark_test["load_kw"].to_frame("actual_load_kw")
```

```
benchmark_test = benchmark_test.drop(columns=["load_kw"])
```

[27]: *# some parts give ConvergenceWarnings here and we'll ignore them.*

```
import warnings
warnings.filterwarnings('ignore')
```

[28]: *# use the same method as before to predict a portfolio for 2020-11-01 -> 2020-11-15*

```
benchmark_test_portfolio = enda.Contracts.forecast_using_trend(
    portfolio_df=portfolio[portfolio.index < '2020-11-01'],
    start_forecast_date=pd.to_datetime("2020-11-01 00:00:00+01:00"),
    nb_days=15,
    past_days=150 # only use recent portfolio trend to forecast the next few days
)
benchmark_test['kva'] = benchmark_test_portfolio['kva']
benchmark_test['contracts_count'] = benchmark_test_portfolio['contracts_count']
benchmark_test
```

[28]:

	contracts_count	kva	tso_forecast_load_mw	\
--	-----------------	-----	----------------------	---

time	contracts_count	kva	tso_forecast_load_mw	\
2020-11-01 00:00:00+01:00	95079.6	808742.3	47900.0	
2020-11-01 00:30:00+01:00	95080.2	808744.8	45800.0	
2020-11-01 01:00:00+01:00	95080.8	808747.2	43700.0	
2020-11-01 01:30:00+01:00	95081.4	808749.7	43900.0	
2020-11-01 02:00:00+01:00	95081.9	808752.2	43200.0	
...	
2020-11-15 21:30:00+01:00	95500.0	810510.5	46200.0	
2020-11-15 22:00:00+01:00	95500.6	810512.9	45200.0	
2020-11-15 22:30:00+01:00	95501.2	810515.4	46400.0	
2020-11-15 23:00:00+01:00	95501.8	810517.9	48600.0	
2020-11-15 23:30:00+01:00	95502.4	810520.4	49400.0	

	t_weighted	t_smooth	minuteofday	dayofweek	\
--	------------	----------	-------------	-----------	---

time	t_weighted	t_smooth	minuteofday	dayofweek	\
2020-11-01 00:00:00+01:00	12.67	12.37	0	6	
2020-11-01 00:30:00+01:00	12.68	12.37	30	6	
2020-11-01 01:00:00+01:00	12.70	12.37	60	6	
2020-11-01 01:30:00+01:00	12.66	12.37	90	6	
2020-11-01 02:00:00+01:00	12.63	12.36	120	6	
...	
2020-11-15 21:30:00+01:00	12.05	12.01	1290	6	
2020-11-15 22:00:00+01:00	11.92	11.97	1320	6	
2020-11-15 22:30:00+01:00	11.84	11.96	1350	6	
2020-11-15 23:00:00+01:00	11.75	11.94	1380	6	
2020-11-15 23:30:00+01:00	11.64	11.92	1410	6	

	month	minuteofday_cos	minuteofday_sin	\
time				
2020-11-01 00:00:00+01:00	11	1.000000	0.000000	
2020-11-01 00:30:00+01:00	11	0.991445	0.130526	
2020-11-01 01:00:00+01:00	11	0.965926	0.258819	
2020-11-01 01:30:00+01:00	11	0.923880	0.382683	
2020-11-01 02:00:00+01:00	11	0.866025	0.500000	
...	
2020-11-15 21:30:00+01:00	11	0.793353	-0.608761	
2020-11-15 22:00:00+01:00	11	0.866025	-0.500000	
2020-11-15 22:30:00+01:00	11	0.923880	-0.382683	
2020-11-15 23:00:00+01:00	11	0.965926	-0.258819	
2020-11-15 23:30:00+01:00	11	0.991445	-0.130526	

	dayofweek_cos	dayofweek_sin	dayofyear_cos	\
time				
2020-11-01 00:00:00+01:00	0.62349	-0.781831	0.500000	
2020-11-01 00:30:00+01:00	0.62349	-0.781831	0.500000	
2020-11-01 01:00:00+01:00	0.62349	-0.781831	0.500000	
2020-11-01 01:30:00+01:00	0.62349	-0.781831	0.500000	
2020-11-01 02:00:00+01:00	0.62349	-0.781831	0.500000	
...	
2020-11-15 21:30:00+01:00	0.62349	-0.781831	0.691771	
2020-11-15 22:00:00+01:00	0.62349	-0.781831	0.691771	
2020-11-15 22:30:00+01:00	0.62349	-0.781831	0.691771	
2020-11-15 23:00:00+01:00	0.62349	-0.781831	0.691771	
2020-11-15 23:30:00+01:00	0.62349	-0.781831	0.691771	

	dayofyear_sin	lockdown	public_holiday	\
time				
2020-11-01 00:00:00+01:00	-0.866025	0.0	1.0	
2020-11-01 00:30:00+01:00	-0.866025	0.0	1.0	
2020-11-01 01:00:00+01:00	-0.866025	0.0	1.0	
2020-11-01 01:30:00+01:00	-0.866025	0.0	1.0	
2020-11-01 02:00:00+01:00	-0.866025	0.0	1.0	
...	
2020-11-15 21:30:00+01:00	-0.722117	0.0	0.0	
2020-11-15 22:00:00+01:00	-0.722117	0.0	0.0	
2020-11-15 22:30:00+01:00	-0.722117	0.0	0.0	
2020-11-15 23:00:00+01:00	-0.722117	0.0	0.0	
2020-11-15 23:30:00+01:00	-0.722117	0.0	0.0	



	nb_school_areas_off	extra_long_weekend
time		
2020-11-01 00:00:00+01:00	3.0	0.0
2020-11-01 00:30:00+01:00	3.0	0.0

2020-11-01 01:00:00+01:00	3.0	0.0
2020-11-01 01:30:00+01:00	3.0	0.0
2020-11-01 02:00:00+01:00	3.0	0.0
...
2020-11-15 21:30:00+01:00	0.0	0.0
2020-11-15 22:00:00+01:00	0.0	0.0
2020-11-15 22:30:00+01:00	0.0	0.0
2020-11-15 23:00:00+01:00	0.0	0.0
2020-11-15 23:30:00+01:00	0.0	0.0

[720 rows x 18 columns]

```
[29]: # compare portfolio forecast to reality
for c in ["contracts_count", "kva"]:
    to_plot = pd.merge(
        portfolio[(portfolio.index >= '2020-08-01') & (portfolio.index <=
        ↪ '2020-11-16')][c].to_frame(c+"_historic"),
        benchmark_test[c].to_frame(c+"_forecast"),
        how='outer', left_index=True, right_index=True
    )

    to_plot.plot(figsize=(16, 4))
```



```
[30]: # Lets define some algorithms then train and predict with them
# All the models we define implement the enda.models.ModelInterface (see the
# docs)
# Enda comes with wrappers around scikit-learn and H2O models:
# sklearn: enda.ml_backends.sklearn_estimator.SklearnEstimator
# H2O: enda.ml_backends.h2o_estimator.H2OEstimator

import time
import h2o
import random
import numpy

from sklearn.linear_model import LinearRegression, SGDRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

from enda.ml_backends.h2o_estimator import H2OEstimator # enda's wrapper
# around H2O models
from h2o.estimators import H2OGeneralizedLinearEstimator
from h2o.estimators import H2OXGBoostEstimator
from h2o.estimators import H2OGradientBoostingEstimator
from h2o.estimators import H2ORandomForestEstimator
from h2o.estimators import H2ODeepLearningEstimator
```

```
[31]: random.seed(17) # set random seed for reproducibility
numpy.random.seed(17) # for sklearn
# for h2o we will define it in each model
```

```
[32]: all_models = dict()
```

```
[33]: # Some models with the sklearn machine learning backend

all_models['sklearn_lin_reg'] = SklearnEstimator(LinearRegression())

all_models['sklearn_sgd'] = SklearnEstimator(
    Pipeline([('standard_scaler', StandardScaler()),
              ('sgd', SGDRegressor())
            ])
)

all_models['sklearn_ada_boost'] = SklearnEstimator(AdaBoostRegressor(
    n_estimators=500,
    loss='linear', # 'square'
    learning_rate=0.8))
```

```

)

all_models['sklearn_nn'] = SklearnEstimator(
    Pipeline([('standard_scaler', StandardScaler()),
               ('mlp', MLPRegressor(
                   solver='adam',
                   activation='relu',
                   hidden_layer_sizes=[48, 48, 24],
                   max_iter=150
               ))
            ])
)

```

[34]: *# Some models with the h2o machine learning backend*

```

all_models['h2o_glm'] = H2OEstimator(H2OGeneralizedLinearEstimator(
    standardize=False,
    intercept=True,
    seed=17)
)

all_models['h2o_rf'] = H2OEstimator(H2ORandomForestEstimator(
    ntrees=300,
    max_depth=15,
    sample_rate=0.8,
    min_rows=10,
    nbins=52,
    mtries=3,
    seed=17
))

all_models['h2o_gbm'] = H2OEstimator(H2OGradientBoostingEstimator(
    ntrees=500,
    max_depth=5,
    sample_rate=0.5,
    min_rows=5,
    seed=17
))

all_models['h2o_xgboost'] = H2OEstimator(H2OXGBoostEstimator(
    **{
        "ntrees": 500,
        "max_depth": 5,
        "sample_rate": 0.8,
        "min_rows": 10,
        "seed": 17
    }
)

```

```

    }
))

all_models['h2o_nn'] = H2OEstimator(H2ODeepLearningEstimator(
    **{
        "activation": "Tanh",
        "hidden": [48, 48, 24],
        "distribution": "gaussian",
        "epochs": 20,
        "seed": 17
    }
))

```

[35]: *# You can add more models to the benchmark here if you like*

[36]: *# to train or predict with H2O models, we boot up a local h2o server*

```

h2o.init(nthreads=-1)
h2o.no_progress()

```

Checking whether there is an H2O instance running at http://localhost:54321
... not found.

Attempting to start a local H2O server...

Java Version: java version "12.0.1" 2019-04-16; Java(TM) SE Runtime Environment (build 12.0.1+12); Java HotSpot(TM) 64-Bit Server VM (build 12.0.1+12, mixed mode, sharing)

Starting server from /Users/emmanuel.charon/Documents/CodeProjects/enercoop/enda/venv/lib/python3.7/site-packages/h2o/backend/bin/h2o.jar

Ice root: /var/folders/5x/409ks2012xxch_pmbs6qpzfh0000gp/T/tmpx3_pia1b

JVM stdout: /var/folders/5x/409ks2012xxch_pmbs6qpzfh0000gp/T/tmpx3_pia1b/h2o_emmanuel_charon_started_from_python.out

JVM stderr: /var/folders/5x/409ks2012xxch_pmbs6qpzfh0000gp/T/tmpx3_pia1b/h2o_emmanuel_charon_started_from_python.err

Server is running at http://127.0.0.1:54321

Connecting to H2O server at http://127.0.0.1:54321 ... successful.

```

-----
↪-----
H2O_cluster_uptime:      02 secs
H2O_cluster_timezone:    Europe/Paris
H2O_data_parsing_timezone: UTC
H2O_cluster_version:     3.32.0.4
H2O_cluster_version_age:  1 month and 23 days
H2O_cluster_name:        H2O_from_python_emmanuel_charon_xqjnib
H2O_cluster_total_nodes: 1
H2O_cluster_free_memory: 4 Gb
H2O_cluster_total_cores: 4
H2O_cluster_allowed_cores: 4
H2O_cluster_status:      accepting new members, healthy

```

```

H2O_connection_url:      http://127.0.0.1:54321
H2O_connection_proxy:    {"http": null, "https": null}
H2O_internal_security:   False
H2O_API_Extensions:      Amazon S3, XGBoost, Algos, AutoML, Core V3,
    ↪ TargetEncoder, Core V4
Python_version:          3.7.6 final
-----
    ↪ -----

```

```

[37]: # this should take between 5 and 15 minutes to run (in function of your
    ↪ hardware)
print("Benchmark with {} models : {}\n".format(len(all_models), list(all_models.
    ↪ keys()))))
for model_name, model in all_models.items():
    model_start_time = time.time()
    print("Training {} before predicting with it..".format(model_name))
    model.train(benchmark_train, target_col='load_kw')
    model_prediction = model.predict(benchmark_test, target_col='load_kw')
    benchmark[model_name] = model_prediction
    print("{} took {:.1f} seconds.\n".format(model_name, time.
    ↪ time()-model_start_time))

```

```

Benchmark with 9 models : ['sklearn_lin_reg', 'sklearn_sgd',
'sklearn_ada_boost', 'sklearn_nn', 'h2o_glm', 'h2o_rf', 'h2o_gbm',
'h2o_xgboost', 'h2o_nn']

```

```

Training sklearn_lin_reg before predicting with it..
sklearn_lin_reg took 0.1 seconds.

```

```

Training sklearn_sgd before predicting with it..
sklearn_sgd took 1.3 seconds.

```

```

Training sklearn_ada_boost before predicting with it..
sklearn_ada_boost took 62.1 seconds.

```

```

Training sklearn_nn before predicting with it..
sklearn_nn took 59.4 seconds.

```

```

Training h2o_glm before predicting with it..
h2o_glm took 5.4 seconds.

```

```

Training h2o_rf before predicting with it..
h2o_rf took 33.2 seconds.

```

```

Training h2o_gbm before predicting with it..
h2o_gbm took 22.7 seconds.

```


Training h2o_xgboost before predicting with it..
h2o_xgboost took 64.1 seconds.

Training h2o_nn before predicting with it..
h2o_nn took 80.2 seconds.

[38]: benchmark




```
[38]:
```




	actual_load_kw	sklearn_lin_reg	sklearn_sgd	\
time				
2020-11-01 00:00:00+01:00	6817.332090	7416.477529	7542.312564	
2020-11-01 00:30:00+01:00	6326.667322	7192.916989	7319.436451	
2020-11-01 01:00:00+01:00	6172.223671	6974.931684	7102.126243	
2020-11-01 01:30:00+01:00	6050.575318	6993.462928	7122.302861	
2020-11-01 02:00:00+01:00	5898.881230	6928.525322	7058.529345	
...	
2020-11-15 21:30:00+01:00	7657.293444	7554.874780	7696.545257	
2020-11-15 22:00:00+01:00	7317.540759	7422.568321	7565.035767	
2020-11-15 22:30:00+01:00	7580.051439	7510.802466	7655.102382	
2020-11-15 23:00:00+01:00	7496.273993	7702.238171	7848.835422	
2020-11-15 23:30:00+01:00	7376.005701	7759.626894	7908.068267	
	sklearn_ada_boost	sklearn_nn	h2o_glm	\
time				
2020-11-01 00:00:00+01:00	6822.188092	7351.172908	7402.432609	
2020-11-01 00:30:00+01:00	6414.531562	7008.480631	7164.175176	
2020-11-01 01:00:00+01:00	6260.174000	6668.162284	6925.916729	
2020-11-01 01:30:00+01:00	6260.174000	6546.138912	6948.635651	
2020-11-01 02:00:00+01:00	6162.218857	6329.058216	6869.233390	
...	
2020-11-15 21:30:00+01:00	7679.641519	8075.749892	7227.449423	
2020-11-15 22:00:00+01:00	7632.927837	7831.622914	7114.005755	
2020-11-15 22:30:00+01:00	7726.430994	7792.580825	7250.192657	
2020-11-15 23:00:00+01:00	7823.412347	7839.531977	7499.847540	
2020-11-15 23:30:00+01:00	7823.412347	7739.108252	7590.647251	
	h2o_rf	h2o_gbm	h2o_xgboost	h2o_nn
time				
2020-11-01 00:00:00+01:00	6537.198906	7010.410149	7032.040527	7400.976455
2020-11-01 00:30:00+01:00	6228.407756	6408.337175	6492.153809	6841.493779
2020-11-01 01:00:00+01:00	6058.203420	6207.136015	6351.727051	6353.977698
2020-11-01 01:30:00+01:00	6019.727900	6059.606349	6259.075684	6178.509022
2020-11-01 02:00:00+01:00	5838.660528	5967.455171	6184.107422	5974.518119
...
2020-11-15 21:30:00+01:00	7483.921108	7272.697440	7436.578125	8017.731406
2020-11-15 22:00:00+01:00	7369.696343	7050.622522	7187.884766	7788.172288




2020-11-15 22:30:00+01:00	7411.223903	7247.823333	7486.856934	7715.353810
2020-11-15 23:00:00+01:00	7425.497777	7294.028407	7286.999512	7702.300630
2020-11-15 23:30:00+01:00	7458.976123	7117.461922	7210.519043	7612.710499

[720 rows x 10 columns]

```
[39]: # visualize predictions
for c in benchmark.columns:
    if c != "actual_load_kw":
        to_plot = benchmark[["actual_load_kw", c]]
        to_plot.plot(ylim=0, figsize=(16, 4))
```







```
[40]: # compute the mean absolute percentage error
benchmark_ape = benchmark.copy(deep=True).drop(columns=["actual_load_kw"])
for c in benchmark_ape.columns:
    benchmark_ape[c] = (benchmark_ape[c] - benchmark["actual_load_kw"]).abs() /
    ↪ benchmark["actual_load_kw"] * 100
benchmark_ape.mean().to_frame("mape")
```

```
[40]:          mape
sklearn_lin_reg    7.055730
```

sklearn_sgd	7.513167
sklearn_ada_boost	6.347616
sklearn_nn	5.091200
h2o_glm	9.057864
h2o_rf	3.424008
h2o_gbm	1.721648
h2o_xgboost	1.995853
h2o_nn	6.338839

1.6 6. Benchmark with Backtesting

In traditional machine learning, we need more than just 1 evaluation to test an algorithm. We typically use cross-validation to see if the algorithm is not biased and if it can be expected to work well in most cases. For time-series predictions we cannot do a regular cross-validation because it is not realistic : we always want to train using historical data that happened before the prediction.

Here we will do **backtesting** week after week. With the given dataset, this means : - for each week w from early 2019 until the end of the dataset : train using data from the beginning of the dataset (early 2015) until a few days before week w, then eval on w. - the first iteration will train an algorithm using data from 2015 to 2018, then eval on the first week of 2019 - the second iteration will train using data from 2015 to a bit before the first week of 2019, then eval on the second week of 2019 - and so on... - keep the predictions of each time-step using this method, from early 2019 to november 2020.

- then compare these predictions to the historic data to evaluate the quality of each algorithm.

This makes most sense if in your production environment, you plan to retrain the algorithm regularly with recent data.

Backtesting can take a significant amount of time. We backtest only 2 linear regressions below in order to have an example that runs fast. Don't hesitate to add other algorithms.

```
[41]: all_models = dict()

all_models['sklearn_lin_reg'] = SklearnEstimator(LinearRegression())

all_models['h2o_glm'] = H2OEstimator(H2OGeneralizedLinearEstimator(standardize=False,
↪ intercept=True))

[42]: start_backtesting_dt = pd.to_datetime('2019-01-01 00:00:00+01:00').
↪ tz_convert('Europe/Paris')
benchmark = historic[historic.index>=start_backtesting_dt]["load_kw"].
↪ to_frame("actual_load_kw")
days_in_each_iteration = 28

for model_name, model in all_models.items():
```

```

count_iterations = 0

model_predictions = []

for train_set, test_set in enda.BackTesting.yield_train_test(
    historic,
    start_eval_datetime=start_backtesting_dt,
    days_between_trains=days_in_each_iteration,
    gap_days_between_train_and_eval=14
):
    count_iterations += 1
    if count_iterations <= 2 or count_iterations % 10 == 0:
        print("Model {}, backtesting iteration {}, train set {}->{}, test_
→set {}->{}\n".format(
            model_name, count_iterations,
            train_set.index.min(), train_set.index.max(),
            test_set.index.min(), test_set.index.max()))

    # featurize
    train_set = featurize(train_set)
    test_set = test_set.drop(columns=["load_kw"])
    test_set = featurize(test_set)

    # use forecast portfolio in test_set
    forecast_portfolio = enda.Contracts.forecast_using_trend(
        portfolio_df=portfolio[portfolio.index<test_set.index.min()],
        start_forecast_date=test_set.index.min(),
        nb_days=days_in_each_iteration,
        past_days=150) # recent portfolio trend

    test_set['kva'] = forecast_portfolio['kva']
    test_set['contracts_count'] = forecast_portfolio['contracts_count']

    # train and predict
    model.train(train_set, target_col='load_kw')
    model_predictions.append(model.predict(test_set, target_col='load_kw'))

benchmark[model_name] = pd.concat(model_predictions)

```

```

Model sklearn_lin_reg, backtesting iteration 1, train set 2015-01-01
00:00:00+01:00->2018-12-17 23:30:00+01:00, test set 2019-01-01
00:00:00+01:00->2019-01-28 23:30:00+01:00

```

```

Model sklearn_lin_reg, backtesting iteration 2, train set 2015-01-01
00:00:00+01:00->2019-01-14 23:30:00+01:00, test set 2019-01-29
00:00:00+01:00->2019-02-25 23:30:00+01:00

```

Model sklearn_lin_reg, backtesting iteration 10, train set 2015-01-01
00:00:00+01:00->2019-08-26 23:30:00+02:00, test set 2019-09-10
00:00:00+02:00->2019-10-07 23:30:00+02:00

Model sklearn_lin_reg, backtesting iteration 20, train set 2015-01-01
00:00:00+01:00->2020-06-01 23:30:00+02:00, test set 2020-06-16
00:00:00+02:00->2020-07-13 23:30:00+02:00


Model h2o_glm, backtesting iteration 1, train set 2015-01-01
00:00:00+01:00->2018-12-17 23:30:00+01:00, test set 2019-01-01
00:00:00+01:00->2019-01-28 23:30:00+01:00


Model h2o_glm, backtesting iteration 2, train set 2015-01-01
00:00:00+01:00->2019-01-14 23:30:00+01:00, test set 2019-01-29
00:00:00+01:00->2019-02-25 23:30:00+01:00

Model h2o_glm, backtesting iteration 10, train set 2015-01-01
00:00:00+01:00->2019-08-26 23:30:00+02:00, test set 2019-09-10
00:00:00+02:00->2019-10-07 23:30:00+02:00

Model h2o_glm, backtesting iteration 20, train set 2015-01-01
00:00:00+01:00->2020-06-01 23:30:00+02:00, test set 2020-06-16
00:00:00+02:00->2020-07-13 23:30:00+02:00

```
[43]: # visualize predictions
for c in benchmark.columns:
    if c != "actual_load_kw":
        to_plot = benchmark[["actual_load_kw", c]]
        to_plot.plot(ylim=0, figsize=(16, 4))
```





```
[44]: # compute absolute percentage error
benchmark_ape = benchmark.copy(deep=True).drop(columns=["actual_load_kw"])
for c in benchmark_ape.columns:
    benchmark_ape[c] = (benchmark_ape[c] - benchmark["actual_load_kw"]).abs() /
    ↪ benchmark["actual_load_kw"] * 100
benchmark_ape.mean().to_frame("mape")
```

```
[44]:
```

	mape
sklearn_lin_reg	13.164438
h2o_glm	14.279924

If you have time/computing power: - try more algorithms in the backtesting benchmark, this is longer but more reliable than a simple benchmark (think of it as crossval versus single eval in a non-time-series setup). - reduce the “days_in_each_iteration” down to 7 if you think you can have a weekly training in your production environment.

1.7 7. Make the prediction

Seeing the results from just the basic benchmark, we here decide to predict using h2o’s gbm. We now need to train it on the full dataset and make the prediction.

```
[45]: gbm = H2OEstimator(H2OGradientBoostingEstimator(
    ntrees=500,
    max_depth=5,
    sample_rate=0.5,
    min_rows=5
))
```

```
[46]: gbm.train(full_train_set, target_col='load_kw')
```

```
[47]: gbm_prediction = gbm.predict(forecast_input_data, target_col="load_kw")
```

```
[48]:
```




```

# visualize recent load along with our forecast; remember we don't have recent
↳ actual load so there is a time-gap.
# (remember that the prediction takes weather forecast and more information
↳ into account)
to_plot = pd.merge(
    historic["load_kw"][-1000:].to_frame("historic_kw"),
    gbm_prediction.rename(columns={"load_kw": "forecast_kw"}),
    how='outer', left_index=True, right_index=True
)
to_plot.plot(ylim=0, figsize=(16, 4))

```

[48]: <AxesSubplot:xlabel='time'>



```

[49]: # don't forget to shutdown your h2o local server
h2o.cluster().shutdown()
# wait for h2o to really finish shutting down
time.sleep(5)

```

H2O session _sid_88cb closed.

1.8 Conclusion

Thats all for Example B. Check out Example C next. Thanks for reading and don't hesitate to send feedback at: emmanuel.charon@enercoop.org !